

A Complete Set of Related Git Repositories Identified via Community Detection Approaches Based on Shared Commits

Audris Mockus

The University of Tennessee
Knoxville, Tennessee
audris@mockus.org

Zoe Kotti

Athens University of Economics and Business
Athens, Greece
zoeckotti@hotmail.com

Diomidis Spinellis

Athens University of Economics and Business
Athens, Greece
dds@aueb.gr

Gabriel John Dusing

The University of Tennessee
Knoxville, Tennessee
gdusing@vols.utk.edu

ABSTRACT

In order to understand the state and evolution of the entirety of open source software we need to get a handle on the set of distinct software projects. Most of open source projects presently utilize Git, which is a distributed version control system allowing easy creation of clones and resulting in numerous repositories that are almost entirely based on some parent repository from which they were cloned. Git commits are unlikely to get produce and represent a way to group cloned repositories. We use World of Code infrastructure containing approximately 2B commits and 100M repositories to create and share such a map. We discover that the largest group contains almost 14M repositories most of which are unrelated to each other. As it turns out, the developers can push git object to an arbitrary repository or pull objects from unrelated repositories, thus linking unrelated repositories. To address this, we apply Louvain community detection algorithm to this very large graph consisting of links between commits and projects. The approach successfully reduces the size of the megacluster with the largest group of highly interconnected projects containing under 400K repositories. We expect that the resulting map of related projects as well as tools and methods to handle the very large graph will serve as a reference set for mining software projects and other applications. Further work is needed to determine different types of relationships among projects induced by shared commits and other relationships, for example, by shared source code or similar filenames.

KEYWORDS

forks and clones

ACM Reference Format:

Audris Mockus, Diomidis Spinellis, Zoe Kotti, and Gabriel John Dusing. 2020. A Complete Set of Related Git Repositories Identified via Community Detection Approaches Based on Shared Commits. In *17th International Conference on Mining Software Repositories (MSR '20)*, October 5–6, 2020,

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

MSR '20, October 5–6, 2020, Seoul, Republic of Korea

© 2020 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-7517-7/20/05...\$15.00

<https://doi.org/10.1145/3379597.3387499>

Seoul, Republic of Korea. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3379597.3387499>

1 INTRODUCTION

While study of individual software projects has been ongoing for some time, relatively less effort has been spent on studying groups of projects even though numerous benefits of understanding groups of projects exist. Extensive motivation for investigating groups of projects was stated in, for example [12, 13]. Furthermore, several attempts at creating an infrastructure for such studies have been reported [8–11]. Here we focus on an aspect of apparently simple, but very hard to address challenge of identifying related repositories. By related repositories we mean repositories that are “developed for the same project/component.” Unrelated repositories are, thus, are not intending to merge their code to a single project. For example, while most GitHub forks are not meant to be independent projects, but some are. Once we go beyond GitHub, often no information about relatedness is available. For, example, searching for relevant code or looking for a project to join, the massive number of the repositories would waste time and cause confusion.

Simply stated, if we obtain data from two distinct git repositories, should we treat it as belonging to a single project or as belonging to two unrelated projects. This is an important question since many projects have tens of thousands of forks that often contain no or very little original content as the fork was created just to submit a pull request and may not have even been used for that purpose. Some relief can be found for projects on GitHub, where forked projects can be identified via GitHub API. No such information is available for other projects and projects that were not forked via GitHub API will be missed. Having a reference list of related projects can provide a massive help to Mining Software Repositories community by providing a common basis that everyone can use to count, identify, sample from, or analyze projects. Our operational definition of a set of repositories representing an independent project is that all these repositories share the objective to work on the same project. For example, a fork created to submit a pull request or to fix a bug that, for all practical purposes, is expected to be eventually fixed upstream can be illustrated by the repositories Debian distribution uses to keep track of upstream packages. They are used to ensure that everything compiles and can be installed together for the distribution, but are not intended to maintain the upstream project. This also includes cases where

a project may maintain another project within its own repository, but not with the purpose of developing it, just to avoid potential incompatibilities that may occur due to differences in development schedule. Hard forks, on the other hand, would indicate a desire to develop the project independently and should be considered as separate projects.

While there are many ways to identify related projects, here we focus on a single approach: linking projects sharing at least one commit. Git commits are based on Merkle Tree and no two commits are likely to be produced independently. For example, the initial commit to a repo creating an empty README.md file and done at exactly the same time (up to a second) in the same timezone, by two developers having identical credentials would result in an identical commit. However, as a distributed VCS, git makes it easy to create clones (via git clone or through GitHub fork button) and resulting in numerous repositories that are intended to be distributed copies of the code used in the same project. This feature of Git that enables distributed collaboration also results in numerous clones of the original repository. Furthermore, GitHub introduced single-click way to fork (in essence to clone) a repository on GitHub and use it to create patches (pull requests) for the origin of that fork. This further increased the number repositories related to popular projects.

As noted above, it is virtually impossible to produce independently identical commits in normal development (see a potential example above), so it would appear that projects sharing a commit are related. Here we do not consider other types of related projects where the version history was not shared and only the source code has been copied. Such projects can not be identified via shared commits and are a subject of further work, for example by comparing the blobs shared among the projects or the directory structure of the source code [6, 7, 19].

To apply the approach we utilize the infrastructure provided by the World of Code [14]. Specifically, we use version Q of the data and obtain commit to project relationships. As described in WoC tutorial [15], the data is stored in 32 databases containing a full list of pairs between commits and projects in which these commits were found (c2pFullQXX.s). This commit to project graph has a total of 99,154,451,345 links between 116265607 projects and 1868632121 commits. We handle the scale of the problem as described in Section 3, i.e., by solving a sequence of smaller problems and using the results to solve the larger problem. The largest group of repositories (we use words “project” and “repository” interchangeably here) has almost 14M projects and not all of them appear to be closely related.

We, therefore identify some of the reasons for such outcome (projects that fetch from or push to repositories of unrelated projects) and propose and implement alternative operationalizations of related projects. These involve the attempt to identify and remove problematic projects or commits, simplification of the problem by reducing the number of projects by using explicit fork identification in GitHub, and using Louvain community detection algorithm to separate connected but unrelated projects.

Our results provide an operationalization of related projects for open source projects utilizing git version control system obtained from WoC infrastructure. In addition to projects related by a shared commit, we also provide the ultimate parent from forking relationships for GitHub forks and the groups defined by Louvain community detection algorithm.

The remainder of the paper describes data sources in Section 2, the approach to link the projects in Section 3, the approach to eliminate problematic projects and commits in Section 4, the community detection approach described in Section 5, and summary in Section 9.

2 DATA SOURCE: WORLD OF CODE

The World of Code [14] infrastructure prototype was created to support developing theoretical, computational, and statistical frameworks that discover, collect, and process FLOSS operational data and construct FLOSS supply chains (SC), identify and quantify its risks, and discover and construct effective risk mitigation practices and tools. That prototype stores the huge and rapidly growing amount of data in the entire FLOSS ecosystem and provide basic capabilities to efficiently extract and analyze that data at that scale. WoC’s primary focus is on types of analyses that require global reach across FLOSS projects.

In a nutshell, WoC is a software analysis pipeline starting from discovery and retrieval of data, storage and updates, and transformations and data augmentation necessary for analytic tasks downstream. In addition to storing objects from all git repositories WoC also provides relationships among them. For the purpose of this analysis we only use a single relationship from WoC: commit to project map that lists all commit project pairs. WoC has two interfaces: one optimized for random access and another for processing the entirety of the collection. We chose the second due to need to obtain the entirety of the commit to project links. WoC splits each relationship into 32 databases. Specifically, the c2p (commit to Project) database is split based on 5 bits of the first byte of commits Sha1. Thus we naturally have 32 smaller datasets to analyze. Randomness of Sha1 ensures that each of these databases represents the entire collection.

WoC data is versioned with the latest version labeled as Q and containing 7204111388 blobs, 1868632121 commits, 7596825726 trees, 16172556 tags, 116265607 projects (distinct repositories), and 38142898 distinct author IDs. WoC has collected that data during November and December of 2019. For more information please consult WoC website [3]. The proposed grouping into the related projects produced 66532614 such clusters with the largest cluster containing 354920 repositories (miranagha/js).

We also use fork parent data obtained from GHTorrent [11] and, for GitHub projects not present in GHTorrent, we retrieved using GitHub GraphQL API [2]. Please note that that GitHub forks may have their own forks. For each project we obtain the ultimate parent: that is if the parent has a parent, we continue until the repository is no longer fork.

3 LINKING PROJECTS BY SHARED COMMITS

As noted above, we distribute the computational load over the 32 databases listing commit/project pairs. Since data in these lists are sorted we simply need to group projects that share the same commit. Commits belonging to a single project can be ignored as they will not provide a link among projects. The result of the first pass over each of the 32 databases is a list of lines each listing two or more projects linked by a commit (for more detail please see ssc-oscar/forks/README.md). We then encode each line representing a group of N projects as $N - 1$ links linking the first project to

the remaining ones. The resulting graph is used to produce cliques (connected components) via C++ Boost library [1]¹. The resulting components from each of the 32 databases are then combined into a single graph and the same library is used to produce the overall components. The largest components are shown in Table 1. Names

Member Count	Name
13,912,612	grr
28,193	rh24/parrot-ruby
17,267	kvignali/arel-lab
16,181	hmagph/ui-
16,170	54/996iCU
10,541	mil/kb
10,218	bloomni/aa
9,911	f0/rkt

Table 1: The largest groups of related repositories

of the clusters are chosen by selecting a repository from the cluster that has the shortest name that and is first in the alphanumeric order. This cluster name is provided as the second column of the provided map, where the first column lists all 116,265,607 projects and the linking produces 61,921,909 distinct clusters unconnected by commits. The first mega-cluster exceeds the next one by almost three magnitudes and is clearly undesirable as it packs more than 10% of all projects and groups together what appear to be rather unrelated projects.

4 REMOVING BAD PROJECTS AND COMMITS

Given less than ideal outcome obtained in Section 3, we have spent some time investigating the reasons behind that outcome. Specifically we identified at least two kind of repositories that give rise to such a mega-cluster. First, it appears that some projects are used in what appears to be simply a backup storage. Since any developer who has a permission to write to a repository can push git objects to it from any unrelated repository, this feature may have been used by some developers to use cloud git version control systems simply to back up their work. Examples of such repositories include “docker-library/commit-warehouse” and “devilln-side/AcerRecovery.” The second class of problematic repositories appears to include repositories that contain version history from multiple independent projects that are used to build a single project, for example, “bloomberg/chromium.bb” that contains commit history from independent projects such as libdrm and FFmpeg. A simple attempt to remove such projects manually did not give great results as after one of such problematic projects was removed, there we hundreds of others that remained leaving the size of the mega-cluster stubbornly high. After eliminating a large set of potentially problematic projects (listed in the code as an associative array bad Projects) and also removing potentially problematic commits (commits that span more than one thousand projects), we still had a formidable mega-cluster containing 9,626,594 projects and 65,591,526 groups of unrelated projects.

5 COMMUNITY DETECTION

Research on large graphs has produced a number of algorithms that detect communities: groups that interact (have more links)

¹Specifically `connected_components` function from “boost/graph/connected_components.hpp”

among themselves than across groups. Such algorithms tend to be much more time consuming than the arguably simplest connected subset detection algorithm we used in Section 3. More importantly, it is not clear how to combine the results from multiple runs of the algorithm on different subsets of commits as we did in that section. We, therefore, tried to simplify the problem in two ways. First, we reduce the number of distinct projects by using information obtained from GitHub fork API and substituting project name by its ultimate parent. Second, we reduce the number of commits by considering all commits touching the same subset of projects as a single commit. That resulted in 141,53,282 groups (hyperlinks) of projects representing a minimum of 13,82,233,820 links involving two projects.

After this preparation the resulting graph was then analyzed using iGraph package in R [16]. It was necessary to read and add links to the graph in chunks in order to avoid creating a long vector (iGraph can not presently handle R’s long vector). Louvain’s algorithm implementation in R was used, specifically, the function “cluster_louvain” [5]. Louvain’s algorithm optimizes modularity. It was chosen for its speed— spectral clustering requires $O(n^3)$ operations while Louvain is nearly linear on the sparse graph we have.

The resulting set of groups (we use groups, components, and clusters interchangeably) appears to be much more reasonable with the three largest groups representing what appear to be legitimately related groups of projects involved in language tutorials (miranagha/js), github.io templates of creating a static github website/personal CV (6101/-) programming assignments (ykgm/R), datasharing templates (jkwonl/test), linux kernel for mobile mods (aosp/oz), bootstrap (UCF/50), configuration files (rdp/a), and spring framework (maiyy/-).

Member Count	Name
354920	miranagha/js
333645	6101/-
241893	ykgm/R
211538	jkwonl/test
179315	aosp/oz
101988	UCF/50
94160	rdp/a
89602	maiyy/-

Table 2: The largest groups of related repositories using Louvain community detection

6 DATA OVERVIEW

We provide “ultimateMap2.s” where the first column is the repo name transformed with the first ‘/’ replaced by ‘_’, and the ‘github.com/’ removed for GitHub repositories. It was produced by assigning the result from running community detection mentioned previously to assign a cluster name which represents an independently developed project to all repositories in World of Code.

We also provide ‘ghForks.gz’ which is produced by identifying, for each forked repository, its ultimate parent. If a parent is a fork itself, find its parent, and so on, until it is not a fork.

7 EVALUATION

To measure the accuracy of the community detection algorithm we rely on the incomplete list of ultimate parent repositories at

the time we performed the clustering calculation. Over the weeks during which we were doing the clustering, we also retrieve information using GitHub API on whether or not the project was a fork and, if so, what was the parent. Over 15M projects from WoC could not be found in the ghTorrent extract we used. Due to throttling of GitHub API the process of obtaining fork parents is very slow. Over the period we did the computation, we were able to retrieve fork parent information for only approximately one million GitHub repositories. By the time of writing we have collected fork information on 1,652,872 repositories that was not available for the community detection analysis described in this section. We used that information to determine if the community detection approach was able to group these repositories to the corresponding ultimate parents in this new extract. Of these, only 32,082 or 1.9% were not placed in the more than one group (represented by the ultimate parent). Of these incorrectly split, most (9,245) were in the octocat/Spoon-Knife, which is a test repository for developers to practice using git. Also, only a tiny percentage of repositories were separated from the main group, for example only one repository was split from spring-projects/spring-boot (see Table 3. As shown

In split	lrgst grp	Parent fork
9245	9222	octocat/Spoon-Knife
2717	2684	rdpeng/ProgrammingAssignment2
1957	1936	rdpeng/ExData_Plotting1
1046	1045	spring-projects/spring-boot

Table 3: Split fork parents with most repositories

in the table the most repositories in forks that were split, occurred in training repositories and with only a few repositories not in the primary group.

We also compare our approach to the competing approach described [18]. Specifically, the competing approach provides groups for 10,649,348 repositories and the set of repositories that we could match was 8157317. Assuming the competing approach as the gold standard, our approach splits 100,300 of the 2,036,117 groups (5%) in the competing approach. Conversely, assuming our approach as the gold standard, the competing approach splits 44,357 out of 2,124,711 (2%) of the groups we detect using our algorithm.

Inspecting the largest discrepancies, our approach produces 629 groups for the torvalds/linux group of the competing approach and competing approach splits aosp/oz group (kernel mobile mods) produced by our algorithm into 1,245 groups. All data and code needed to reproduce the results are in github.com/ssc-oscar/forks.

8 LIMITATIONS

It is important to note that we are not trying to solve the problem of identifying all related project, just ones that are related via code commits. We are also not investigating finer types of relationships, for example, light forks done for a single pull request vs hard forks where projects evolve independently or all shades of grey in between. Other approaches may be more suitable (or used in combination with shared commit methods) for that. For example, shared code, amount of independent evolution, etc.

We utilize WoC data collection with all associated limitations of using that repository and described there [14].

The accuracy of our approach is not easy to establish. While we rely on explicitly specified GitHub forks in the community

detection step, these may involve cases where the forked projects are developed independently with no intention to merge. We were able to reproduce the explicitly specified forks with high accuracy, however.

Some of the projects in WoC may have been renamed and the new project may have the name of the old project but an entirely different content. Identifying and eliminating such projects would help improve the accuracy of the community detection algorithms.

While our approach appears to provide sensible groups of projects, it may be further improved by experimenting with different community detection algorithms and by weighting links in a different manner.

While there is no particular reason to expect that modularity-focused algorithms should work on a problem involving the discrimination of forks and non-forks, it is not unreasonable to assume that unusual patterns of using git such that objects from unrelated projects are shared within a single repository, would appear as anomalies and thus be eliminated as spurious links between groups of legitimate forks.

9 SUMMARY

The main purpose of this work is to demonstrate the feasibility of solving the problem of finding groups of repositories representing independent projects on a global scale with a high scale of automation, and to share the resulting dataset with the research community for further improvement. We also hope that the resulting map will be incorporated into WoC and other infrastructures such as BoA [10] and SoftwareHeritage [9] to further simplify sampling, counting, and statistical analysis of the open source projects.

Specifically, we discover that a direct application of commit-sharing resulted in the largest group containing almost 14M repositories. This happened because the developers can push git objects to an arbitrary repository and pull objects from unrelated repositories into their repository, thus linking unrelated repositories. We attempted to eliminate such problematic repositories with limited success until we applied Louvain community detection algorithm. The approach successfully reduces the size of the mega-cluster with the two largest groups of highly interconnected projects containing approximately 100K repositories that all appear to be closely related. As future work, it might be worth considering ways to apply time-series methods to observe graph behavior over time. In [17], a prediction framework is presented for certain graph parameters, e.g. modularity or average degree. Reducing the size of the graph in question may also be helpful which [4] suggests how this might be done. Underpinning this approach is that not all edges of the graph are necessary to draw conclusions, and by embedding the graph in a metric space, certain edges close together in some sense can be treated as a single edge.

We expect the tools that the resulting map of related projects as well as tools and methods to handle the very large graph will serve as a reference set for mining software projects and other applications. Further work, however, will be required to determine the different types of relationships among projects induced by shared commits and other relationships, for example, by shared source code or similar filenames.

The work has been partially supported by the following NSF awards: CNS-1925615, IIS-1633437, and IIS-1901102

REFERENCES

- [1] Boost: C++ libraries.
- [2] Github graphql api.
- [3] Oscar - open source supply chains and avoidance of risk.
- [4] Faisal N. Abu-Khzam and Rana H. Mouawi. Concise fuzzy representation of big graphs: a dimensionality reduction approach, 2018.
- [5] J-L Blondel, R Lambiotte Guillaume, and E Lefebvre. Fast unfolding of community hierarchies in large networks. Technical report, arXiv, 1010. <http://arxiv.org/abs/arXiv:0803.0476>.
- [6] Hung-Fu Chang and Audris Mockus. Constructing universal version history. In *ICSE'06 Workshop on Mining Software Repositories*, pages 76–79, Shanghai, China, May 22–23 2006.
- [7] Hung-Fu Chang and Audris Mockus. Evaluation of source code copy detection methods on FreeBSD. In *5th Working Conference on Mining Software Repositories*. ACM Press, May 10–11 2008.
- [8] Jacek Czerwinka, Nachiappan Nagappan, Wolfram Schulte, and Brendan Murphy. Codemine: Building a software development data analytics platform at microsoft. *IEEE software*, 30(4):64–71, 2013.
- [9] Roberto Di Cosmo and Stefano Zacchiroli. Software heritage: Whyandhowtopreservesoftwaresourcecode. ipres 2017, 2017.
- [10] Robert Dyer, Hoan Anh Nguyen, Hridesh Rajan, and Tien N. Nguyen. Boa: A language and infrastructure for analyzing ultra-large-scale software repositories. In *Proceedings of the 35th International Conference on Software Engineering*, ICSE'13, pages 422–431, 2013.
- [11] Georgios Gousios and Diomidis Spinellis. Ghtorrent: Github's data from a firehose. In *Mining software repositories (msr), 2012 9th ieee working conference on*, pages 12–21. IEEE, 2012.
- [12] Randy Hackbarth, Audris Mockus, John Palframan, and David Weiss. Assessing the state of software in a large enterprise. *Journal of Empirical Software Engineering*, 10(3):219–249, 2010.
- [13] Randy Hackbarth, Audris Mockus, John Palframan, and David Weiss. Assessing the state of software in a large enterprise: A 12-year retrospective. In *The Art and Science of Analyzing Software Data*, pages 411–451. Elsevier, 2016.
- [14] Yuxing Ma, Chris Bogart, Sadika Amreen, Russell Zaretzki, and Audris Mockus. World of code: An infrastructure for mining the universe of open source vcs data. In *IEEE Working Conference on Mining Software Repositories*, May 26 2019.
- [15] Audris Mockus. Woc tutorial, 2019.
- [16] R Development Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2008. ISBN 3-900051-07-0.
- [17] Sandipan Sikdar, Niloy Ganguly, and Animesh Mukherjee. Time series analysis of temporal networks. *The European Physical Journal B*, 89(1), Jan 2016.
- [18] Diomidis Spinellis, Zoe Kotti, and Audris Mockus. A dataset for github repository deduplication. *arXiv:2002.02314*, 2020.
- [19] Jiaxin Zhu, Minghui Zhou, and Audris Mockus. The relationship between folder use and the number of forks: A case study on github repositories. In *ESEM*, Torino, Italy, September 2014.